

company_X_employee_attrition

April 20, 2020

```
[46]: #import required packages
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
from sklearn.feature_selection import RFE
from sklearn.linear_model import LogisticRegression

#load data
#Existing employee data in sheet1 into data1
data1 = pd.read_excel('Hash-Analytic-Python-Analytics-Problem-case-study-1.
↳xlsx',sheet_name=1)
#employees who left into data2
data2 = pd.read_excel('Hash-Analytic-Python-Analytics-Problem-case-study-1.
↳xlsx',sheet_name=2)
#introduce column 'Attrition' to indicate attrition if the employees left or
↳stayed
data1.insert(1, 'Attrition', 0)

data2.insert(1, 'Attrition', 1)
```

```
[47]: #combine the two dataframes
full_data = pd.concat([data1,data2])
#the a look at the first few records of the data
full_data.head(5)
```

```
[47]:
```

	Emp ID	Attrition	satisfaction_level	last_evaluation	number_project	\
0	2001	0	0.58	0.74	4	
1	2002	0	0.82	0.67	2	
2	2003	0	0.45	0.69	5	
3	2004	0	0.78	0.82	5	
4	2005	0	0.49	0.60	3	

	average_monthly_hours	time_spend_company	Work_accident	\
0	215	3	0	
1	202	3	0	
2	193	3	0	
3	247	3	0	
4	214	2	0	

	promotion_last_5years	dept	salary
0	0	sales	low
1	0	sales	low
2	0	sales	low
3	0	sales	low
4	0	sales	low

```
[48]: #peak at the last few records of data
full_data.tail(5)
```

```
[48]:      Emp ID  Attrition  satisfaction_level  last_evaluation  number_project  \
3566   14995         1             0.40             0.57             2
3567   14996         1             0.37             0.48             2
3568   14997         1             0.37             0.53             2
3569   14998         1             0.11             0.96             6
3570   14999         1             0.37             0.52             2
```

	average_monthly_hours	time_spend_company	Work_accident	\
3566	151	3	0	
3567	160	3	0	
3568	143	3	0	
3569	280	4	0	
3570	158	3	0	

	promotion_last_5years	dept	salary
3566	0	support	low
3567	0	support	low
3568	0	support	low
3569	0	support	low
3570	0	support	low

```
[49]: #check for missing values
full_data.isnull().any()
```

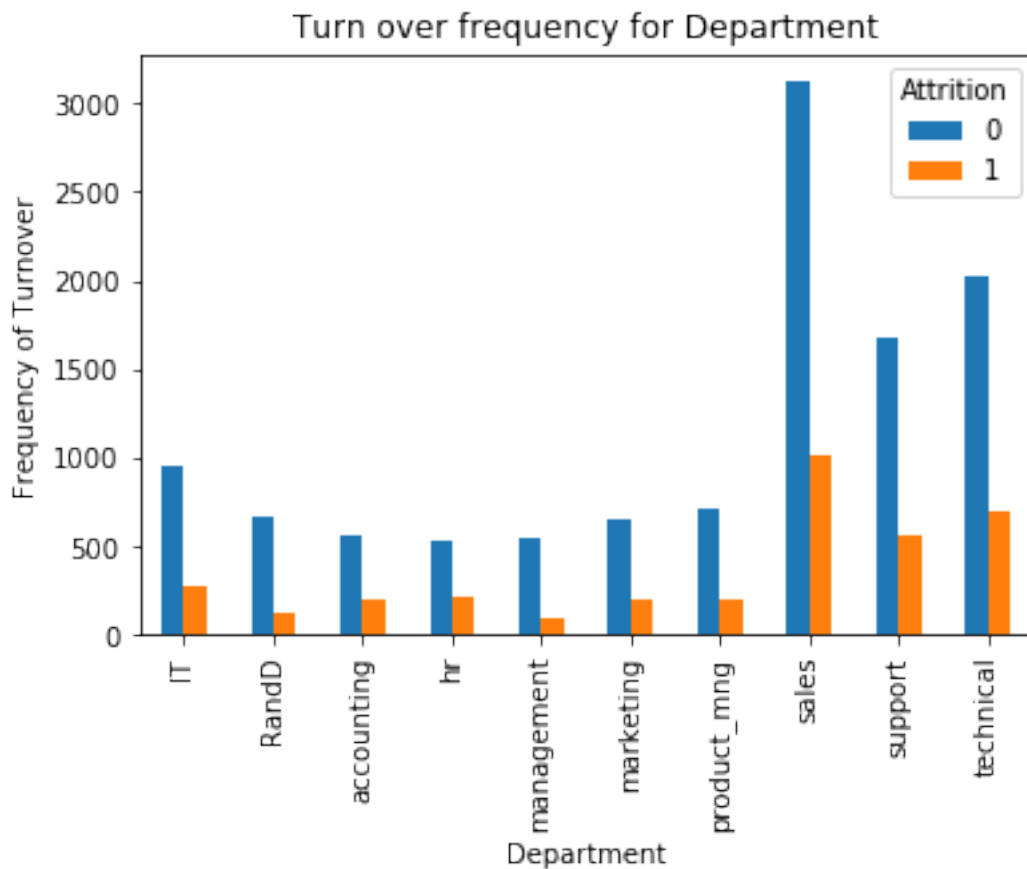
```
[49]: Emp ID           False
Attrition          False
satisfaction_level  False
last_evaluation     False
number_project      False
average_monthly_hours  False
```

```
time_spend_company      False
Work_accident           False
promotion_last_5years   False
dept                   False
salary                 False
dtype: bool
```


```
[50]: #check the number who left and those who did not
full_data["Attrition"].value_counts()
```

```
[50]: 0    11428
      1     3571
      Name: Attrition, dtype: int64
```

```
[51]: #do visualitions
pd.crosstab(full_data.dept,full_data.Attrition).plot(kind='bar')
plt.title('Turn over frequency for Department')
plt.xlabel('Department')
plt.ylabel('Frequency of Turnover')
plt.savefig('department_bar_chart')
```



```
[52]: #Bar chart for employee salary level and the frequency of turnover
table=pd.crosstab(full_data.salary, full_data.Attrition)
table.div(table.sum(1).astype(float), axis=0).plot(kind='bar', stacked=True)
plt.title('Stacked Bar Chart of Salary Level vs Turnover')
plt.xlabel('Salary Level')
plt.ylabel('Proportion of Employees')
plt.savefig('salary_bar_chart')
```

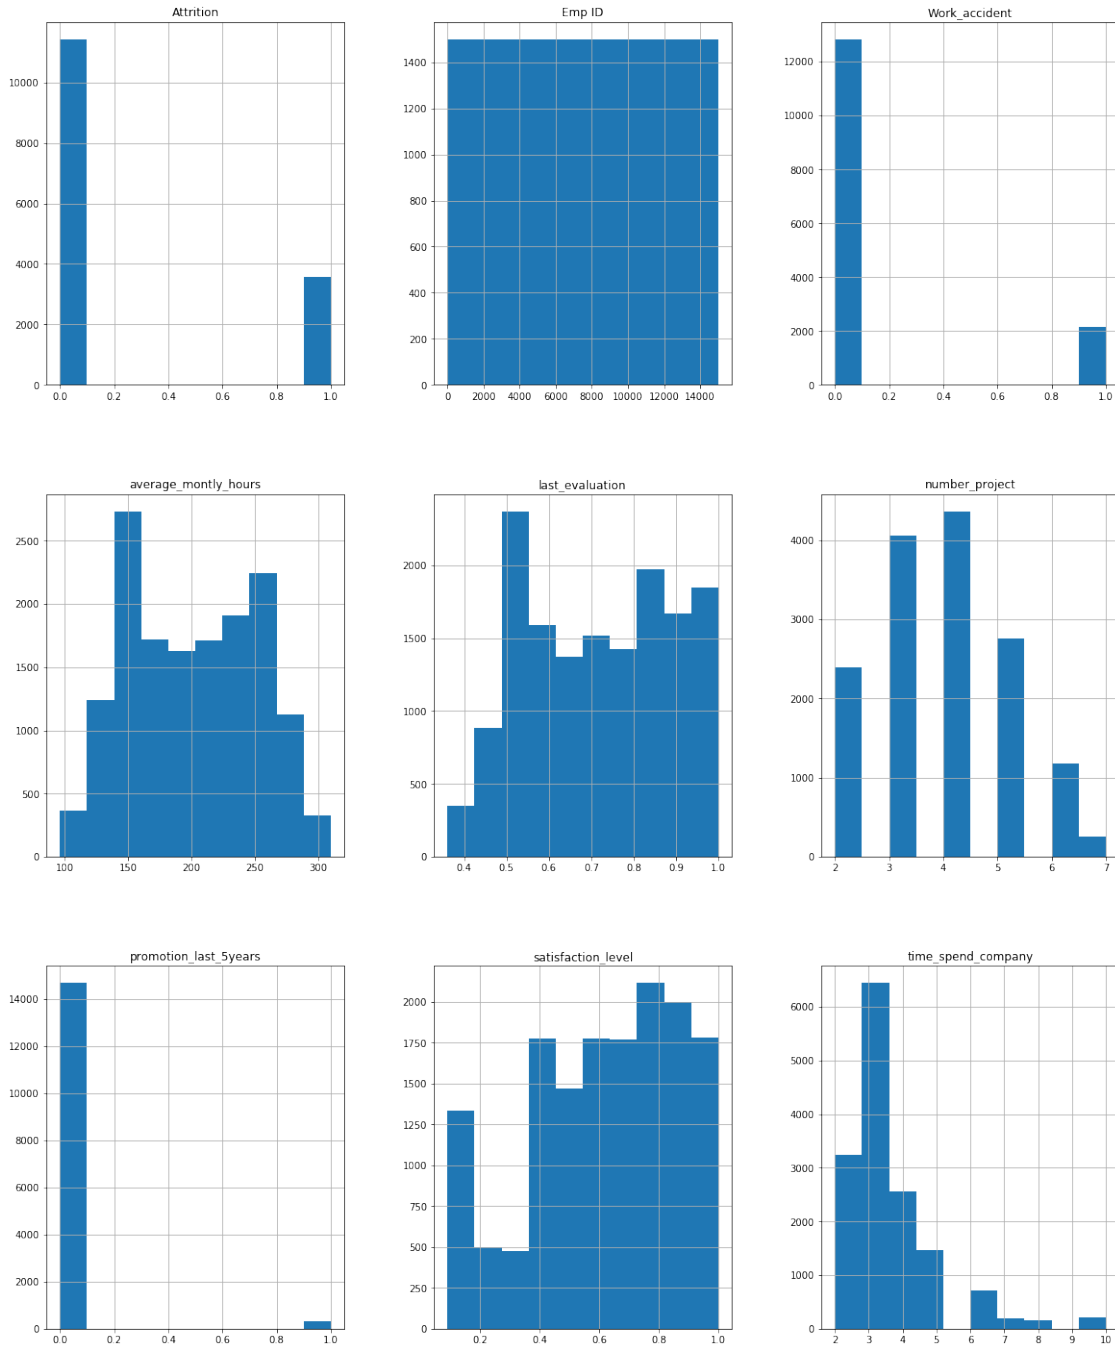


```
[53]: #Proportion of employees who left by department
pd.crosstab(full_data.dept,full_data.Attrition)
```

```
[53]: Attrition      0      1
dept
IT           954    273
RandD        666    121
accounting    563    204
hr           524    215
management    539     91
marketing     655    203
product_mng   704    198
```

sales	3126	1014
support	1674	555
technical	2023	697

```
[54]: #Histogram of numerical variables
num_bins = 10
full_data.hist(bins=num_bins,figsize=(20,25))
plt.savefig("hr_histogram_plots")
plt.show()
```



```
[55]: #create dummy variables for categorical variables
categorical_variables = ['dept','salary']
for var in categorical_variables:
    cat_list = 'var' + '_' + var
    cat_list = pd.get_dummies(full_data[var],prefix=var)
    full_data1 = full_data.join(cat_list)
    full_data = full_data
```

```
[56]: full_data.drop(full_data.columns[[9,10]],axis=1,inplace=True)
```

```
[57]: full_data.columns.values
```

```
[57]: array(['Emp ID', 'Attrition', 'satisfaction_level', 'last_evaluation',
        'number_project', 'average_monthly_hours', 'time_spend_company',
        'Work_accident', 'promotion_last_5years'], dtype=object)
```

```
[58]: full_data_vars = full_data.columns.values.tolist()
y=['Attrition']
X=[i for i in full_data_vars if i not in y]
X
```

```
[58]: ['Emp ID',
        'satisfaction_level',
        'last_evaluation',
        'number_project',
        'average_monthly_hours',
        'time_spend_company',
        'Work_accident',
        'promotion_last_5years']
```

```
[59]: from sklearn.feature_selection import RFE
from sklearn.linear_model import LogisticRegression

#Recursive feature elimination (RFE)
model = LogisticRegression()
rfe = RFE(model,10)
rfe = rfe.fit(full_data[X],full_data[y])
print(rfe.support_)
print(rfe.ranking_)
```

/home/none/.local/lib/python3.8/site-packages/sklearn/utils/validation.py:760: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```

y = column_or_1d(y, warn=True)
[ True  True  True  True  True  True  True  True]
[1 1 1 1 1 1 1 1]

/home/none/.local/lib/python3.8/site-
packages/sklearn/linear_model/_logistic.py:938: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_i = _check_optimize_result(

```

[60]: cols = [ 'satisfaction_level',
               'last_evaluation',
               'number_project',
               'average_monthly_hours',
               'time_spend_company',
               'Work_accident',
               'promotion_last_5years']
X = full_data[cols]
y = full_data['Attrition']

```

```

[65]: #Logistic regression model
      #split the data into training and test samples
      from sklearn.model_selection import train_test_split
      X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.
      ↪3,random_state=0)

```

```

[69]: from sklearn.ensemble import RandomForestClassifier
      from sklearn.metrics import accuracy_score
      #create a model and train
      model = RandomForestClassifier()
      model.fit(X_train,y_train)

```

```


[69]: RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
                             criterion='gini', max_depth=None, max_features='auto',
                             max_leaf_nodes=None, max_samples=None,
                             min_impurity_decrease=0.0, min_impurity_split=None,
                             min_samples_leaf=1, min_samples_split=2,
                             min_weight_fraction_leaf=0.0, n_estimators=100,
                             n_jobs=None, oob_score=False, random_state=None,
                             verbose=0, warm_start=False)

```

```
[70]: #predict the results for the test  
test_pred = model.predict(X_test)  
  
#test the accuracy  
accuracy_score(y_test,test_pred)
```

[70]: 0.9915555555555555

```
[72]: feat_importances = pd.Series(model.feature_importances_,index=X.columns)  
feat_importances= feat_importances.nlargest(20)  
feat_importances.plot(kind='barh')  
plt.savefig('feat_importances_barh')  
plt.show()
```



[]: